

# PRANAV DULEPET

+1(925) 997-0461 ◊ San Ramon, CA

[ps.dulepet@gmail.com](mailto:ps.dulepet@gmail.com) ◊ [linkedin.com/in/pranavdulepet](https://www.linkedin.com/in/pranavdulepet) ◊ [pranavdulepet.github.io](https://github.com/pranavdulepet)

## EDUCATION

<b>M.S.E. Computer Science — Human Language Technology, <i>Johns Hopkins University</i></b>	Aug 2026
<i>Relevant Coursework:</i> Natural Language Processing, Human Language Technology, AI Safety/Alignment/Governance	
<b>B.S. Computer Science — ML, <i>University of Maryland, College Park</i></b>	May 2025
<i>Honors:</i> Dean's List; QUEST Consulting Honors Program	
<i>Relevant Coursework:</i> Systems for Machine Learning, Long-Context LLMs, LLM Privacy & Security, Deep Learning, Machine Learning, Artificial Intelligence, Algorithms, Data Structures	

## SKILLS

**Languages/Technologies** Python, Java, Swift, JavaScript, Git, AWS, closed & open source LLMs

**Libraries/Frameworks** Triton, TensorFlow, PyTorch, Keras, FastAPI, Pandas, MongoDB, Firebase, React

## EXPERIENCE

<b>AI Research Intern, <i>DataTecnica</i></b>	Feb 2026 - Present
• Building AI Agents for biomedical research	
• Creating a novel multi-turn benchmark for the biomedical domain	
• Distilling "aggregation with thinking" into an open model that generalizes to new tools	
<b>Research Assistant, <i>Johns Hopkins CLSP (PI: Benjamin Van Durme)</i></b>	Aug 2025 - Present
• Applying confidence calibration to LLM-based search agents	
<b>Software Engineer Intern (AI/ML), <i>Apple</i></b>	May 2025 - Aug 2025
• AI/LLM infrastructure and API work for Apple Intelligence, all work is under NDA	
• Built prototypes and presented to various teams, and senior leadership	
<b>Technology Policy Researcher, <i>Paragon Policy Fellowship</i></b>	Sep 2024 - May 2025
• Designed a streamlined AI model approval process for the Santa Clara County Government, reducing redundant labor by consolidating three separate interviews into a centralized questionnaire	
• Developed and implemented an AI Usage Guidelines document to assist clients in accurately and comprehensively submitting GenAI applications. <a href="#">View here</a> , <a href="#">AI Use Case Guidelines for Mitigating AI Risk</a>	
<b>Software Developer Intern (AI/ML), <i>Amazon</i></b>	Jun 2024 - Aug 2024
• Developed end-to-end recommendation pipeline using customer-Alexa interaction data (Alexa+ feature)	
• Built data pre-processing framework with PySpark for over 60TB of interaction data	
• Used Claude 3 Sonnet through AWS Bedrock to generate structured outputs from customer interaction data with an acceptance rate of 94%	
<b>Undergraduate Researcher, <i>UMD PIRL (PI: Ramani Duraiswami)</i></b>	Jan 2023 - Jan 2025
• Helped develop a factorable attention mechanism reducing transformers' complexity to $O(N)$ , inspired by fast multipole and Gauss transform methods. <a href="#">View here</a> , <a href="#">FAST</a>	
• Ensured this still captures complete data relationships, avoiding data loss often seen with similar methods	
• Contributed to <a href="#">The Prompt Report</a> , particularly prompt engineering definitions and multimodal sections	
<b>Machine Learning Intern, <i>Capital One</i></b>	Jan 2023 - May 2023
• Implemented <a href="#">NMSLIB similarity search frameworks</a> on financial graph embeddings as part of the Enterprise Graph Services Team to detect transaction fraud	
• Applied to samples of up to 5 million in size with high-dimensional outputting >90 recall (success rate)	

## PROJECTS

**agora.** *Large Language Models, LangChain, Python, SwiftUI, Swift, AWS, MongoDB, Rest APIs*

Developed iOS app and agentic LLM pipeline to provide personalized and affordable meals for students. Adapted Stable Diffusion to generate visuals. Integrated Amazon Fresh and Kroger API to buy ingredients. ([website link](#))

**College RO** *Swift, SwiftUI, Python, Node.js, Rest APIs, MongoDB, AWS, Google/Firebase Analytics*

Launched CollegeRO on the App Store helping college students find research opportunities, reaching a peak of 1.5k app units. ([app link](#))